

# **Exploiting Epochs and Symmetries in Analysing MPI Programs**

**Rishabh Ranjan, Ishita Agrawal, Subodh Sharma**  
IIT Delhi, India

**Automated Software Engineering, 12 Oct 2022, Rochester, MI, USA**

# Introduction

# Motivation

## Verify message-passing (MP) programs for deadlocks

- MP programs are prevalent in high-performance scientific computing, etc.
- The curse of deadlocks:  
easy to introduce, hard to detect, catastrophic manifestations
- Non-determinism makes verification hard:  
undecidable in general, NP-complete for terminating programs
- Message Passing Interface (MPI): a de facto standard for C/C++

# Example - I

$P_0$	$P_1$	$P_2$
$S_{0,0}(1)$	$R_{1,0}(*)$	$S_{2,0}(1)$
$B_{0,1}(0)$	$R_{1,1}(*)$	$B_{2,1}(0)$
$S_{0,2}(1)$	$W_{1,2}(h_{1,1})$	
	$B_{1,3}(0)$	
	$R_{1,4}(*)$	

**Figure 1: Example Message Passing Program**

- $P_i$  → process with rank  $i$   
 $S_{i,k}(j)$  → non-blocking send from  $P_i$  to  $P_j$  at index  $k$   
 $R_{i,k}(j)$  → non-blocking receive from  $P_j$  to  $P_i$  at index  $k$   
 $R_{i,k}(*)$  → wildcard receive to  $P_i$  at index  $k$   
 $W_{i,k}(h_{i,j})$  → blocking wait at index  $k$  for action at index  $j$  for process  $P_i$   
 $B_{i,j}(k)$  →  $k^{\text{th}}$  barrier action from  $P_i$  at index  $j$

# Example - II

$P_0$	$P_1$	$P_2$
$S_{0,0}(1)$	$R_{1,0}(*)$	$S_{2,0}(1)$
$B_{0,1}(0)$	$R_{1,1}(*)$	$B_{2,1}(0)$
$S_{0,2}(1)$	$W_{1,2}(h_{1,1})$	
	$B_{1,3}(0)$	
	$R_{1,4}(*)$	

**Figure 1: Example Message Passing Program**

- **Trace:** sequence of matches allowed by MPI semantics
- **A complete trace:**  $\tau_1 = \langle \{S_{0,0}, R_{1,0}\}, \{S_{2,0}, R_{1,0}\}, \{W_{1,2}\}, \{B_{0,1}, B_{1,3}, B_{2,1}\}, \{S_{0,2}, S_{1,4}\} \rangle$
- **A deadlocking trace:**  
if  $R_{1,1}(*)$  was  $R_{1,1}(0)$ , then  $\tau' = \langle \{S_{0,0}, R_{1,0}\} \rangle$  is a deadlocking trace
- MPI semantics are encoded in:  
a set of allowed matches  $\mathbb{M}$ , and a matches-before order  $<_{mo}$

# Methodology


# Motivation - I

- Communication structure of real-world MPI programs:
  - can be decomposed into **independently-verifiable** epochs
  - has symmetries which lead to **redundancies in the search space** of traces
- Epochs can uncover local symmetries
- Redundant verification of repeated epochs can be avoided
- Symmetry breaking predicates can speed up the search

# Motivation - II

$P_0$	$P_1$	$P_2$
$S_{0,0}(1)$	$R_{1,0}(*)$	$S_{2,0}(1)$
$B_{0,1}(0)$	$R_{1,1}(*)$	$B_{2,1}(0)$
$S_{0,2}(1)$	$W_{1,2}(h_{1,1})$	
	$B_{1,3}(0)$	
	$R_{1,4}(*)$	

**Figure 1: Example Message Passing Program**



**Figure 2: Program Graph for Running Example**



# Epoch Decomposition




Figure 2: Program Graph for Running Example

- **Program Graph (PG):**  
nodes for actions,  
uni-directional edges (solid) for matches-before order,  
bi-directional edge (dashed) for potential matches
- **Epoch:** strongly connected component (SCC) of PG
- **Correctness:** *THEOREM 4.3.  $P$  has a deadlocking trace  $\tau$  if and only if some communication epoch  $e \in E$  has a deadlocking trace  $\tau_e$ .*

# Matchset refinement

- **Potential matches:** allowed matches which can be realized in some run
- Set of potential matches  $\mathbb{M}^+$  is **NP-hard** to compute
- Correctness holds for any **over-approximation**  $M^+$  of  $\mathbb{M}^+$
- But efficacy relies on **tightness** of  $\mathbb{M}^+$
- Start with  $M^+ = \mathbb{M}$  (allowed matches), **refine** using pruning heuristics

# Pruning Heuristics

- Recursive matches-before order pruning:  
for  $c_1$  to match  $c_2$ , all ancestors of  $c_1$  should find a match not successor of  $c_2$
- Barrier-led pruning:  
pairs separated by a barrier cannot match
- Counting heuristic:  
if a set of Sends can match only a set of Recvs of same size,  
then this set of Recvs cannot match other Sends

# Caching

- An epoch once verified is cached
- Epochs isomorphic to cached epochs are skipped
- BLISS package is used for graph isomorphism tests


# Symmetry Detection

- Symmetry is captured by automorphisms of the epoch graph

- Example automorphism for epoch  $E_0$ :

$$\begin{aligned} \pi: S_{0,0}(1) &\mapsto S_{2,0}(1), \\ S_{2,0}(1) &\mapsto S_{0,0}(1), \\ R_{1,0}(\ast) &\mapsto R_{1,0}(\ast), \\ R_{1,1}(\ast) &\mapsto R_{1,1}(\ast). \end{aligned}$$

- Correctness: Traces equivalent under an automorphism are equivalent w.r.t. deadlock



**Figure 2: Program Graph for Running Example**

# Propositional Encoding

- $\phi_t(\tau) \rightarrow \tau$  is a valid trace
- $\phi_d(\tau) \rightarrow \tau$  is a deadlocking trace
- $\phi' = \phi_t \wedge \phi_d$  is SAT iff a deadlocking trace exists
- $\phi_s(\tau) \rightarrow$  satisfied by at least one representative from each equivalence class
- $\phi = \phi_t \wedge \phi_d \wedge \phi_s$  is SAT iff a deadlocking trace exists
- But  $\phi$  is faster to solve, as symmetric branches are pruned from search space

# SAT encoding

- The encoding for  $\phi_t$  and  $\phi_d$  are borrowed from prior work (Mopper)
- $\phi_s$  is encoded using Lex-Leader constraints
- Let  $[\tau]$  denote a bit-vector for trace  $\tau$  in the SAT formula
- For set  $B$  of generators of the automorphism group,  
$$\phi_s = \bigwedge_{\pi \in B} [\tau] \leq [\pi(\tau)],$$
 where  $\leq$  is a lexicographic order
- $B$  is obtained using the BLISS package

# Results



# Implementation

- A prototype tool called **Simian** (<https://github.com/rishabh-ranjan/simian>)
- Components:
  - **Scheduler:** dynamic execution engine, borrowed from prior work (ISP)
  - **Analyzer:** epoch decomposition + symmetry breaking
  - **Solver:** verification of SAT formulas, we use the Z3 solver

# Baselines

- **Mopper:** a SAT encoding for the entire program
- **Mopper-Opt:** alternative SAT encoding optimized for consecutive wildcards
- **Hermes:** a dynamic-symbolic hybrid verifier with an SMT formulation

# Benchmarks

- **Adder:** adds an array of numbers, *master-worker* communication
- **Floyd:** all-pairs shortest path algorithm, *pipelined* communication
- **GaussElim:** Gauss-Jordan elimination, *pairwise* communication
- **Heat:** heat conduction simulation, *2D grid* topology
- **HeatErrors:** Heat, seeded with a *deadlock*
- **Integrate:** numerical integration using trapezoidal rule, *master-worker* communication
- **Diffusion:** iterative solver for diffusion, *2D grid* topology
- **MatMul:** implementation of matrix multiplication, *block-distributed communication over rows*

# Variations

- **Diffusion** with fixed timesteps, but varying number of processes
- **Diffusion** on 2x2 grid, 4 processes, with varying timesteps
- **MatMul** on 8x8 matrices with varying number of processes
- **MatMul** with  $p \times p$  matrices for  $p$  processes
- **MatMul** with fixed number of processes, but varying matrix size

**Table 1: Runtimes for  $\infty$ -buffering (in s)**

Name vs X	X	Deadlock	Mopper	Mopper-Opt	Hermes	Simian
Adder	8	No	0.268	<b>0.042</b>	0.212	0.061
vs	16	No	TO	<b>0.212</b>	TO	1.157
Processes	32	No	TO	<b>1.497</b>	TO	1.912
	64	No	TO	<b>4.116</b>	TO	7.257
Floyd	8	No	2.761	6.132	1.453	<b>0.628</b>
vs	16	No	283.524	399.156	2.148	<b>1.939</b>
Processes	32	No	TO	TO	5.021	<b>2.491</b>
	64	No	TO	TO	10.312	<b>6.081</b>
GaussElim	8	No	0.243	0.233	0.187	<b>0.186</b>
vs	16	No	0.628	1.655	0.258	<b>0.283</b>
Processes	32	No	4.314	4.282	1.993	<b>2.334</b>
	64	No	10.033	6.159	3.912	<b>3.226</b>
Heat	8	No	0.666	0.395	0.406	<b>0.325</b>
vs	16	No	1.581	0.845	0.636	1.506
Processes	32	No	6.543	1.623	2.005	4.255
	64	No	14.927	10.597	4.464	<b>3.232</b>
HeatErrors	8	Yes	0.523	0.395	<b>0.309</b>	0.353
vs	16	Yes	1.191	0.779	0.662	<b>0.565</b>
Processes	32	Yes	5.706	2.574	2.712	<b>2.191</b>
	64	Yes	10.392	<b>4.799</b>	6.435	5.051
Integrate	8	No	0.256	<b>0.038</b>	0.209	0.062
vs	16	No	TO	0.232	TO	<b>0.131</b>
Processes	32	No	TO	3.581	TO	<b>1.854</b>
	64	No	TO	<b>4.212</b>	TO	7.583
Diffusion (Timesteps=1)	4	No	TO	3.729	239.77	<b>0.122</b>
vs	8	No	TO	14.854	TO	<b>0.716</b>
Processes	16	No	TO	49.069	TO	<b>8.566</b>
	24	No	TO	271.247	TO	<b>36.308</b>
Diffusion (Grid=2x2)	2	No	TO	1.598	TO	<b>0.225</b>
vs	4	No	TO	26.347	TO	<b>1.023</b>
Timesteps	8	No	TO	842.951	TO	7.197
	16	No	TO	TO	TO	<b>67.854</b>
MatMul (N=L=M=8)	8	No	2.719	0.083	1.815	<b>0.076</b>
vs	16	No	3.572	0.274	3.032	<b>0.114</b>
Processes	32	No	4.448	<b>0.734</b>	4.014	1.477
	64	No	8.625	4.001	5.482	<b>2.766</b>
MatMul (N=L=M=p)	8	No	2.708	0.086	1.918	<b>0.076</b>
vs	16	No	TO	0.328	TO	<b>0.224</b>
Processes	32	No	TO	3.728	TO	<b>2.822</b>
	64	No	TO	<b>4.808</b>	TO	16.175
MatMul (p=8)	4	No	0.061	<b>0.062</b>	0.071	0.065
vs	6	No	0.101	<b>0.066</b>	0.112	0.079
Size (N=L=M)	8	No	2.704	0.092	1.794	<b>0.076</b>
	12	No	TO	9.772	TO	<b>4.938</b>

**Table 2: Runtimes for 0-buffering (in s)**

Name vs X	X	Deadlock	Mopper	Mopper-Opt	Hermes	Simian
Adder	8	No	0.281	<b>0.043</b>	0.285	0.061
vs	16	No	TO	<b>1.249</b>	TO	1.301
Processes	32	No	TO	1.684	TO	<b>0.932</b>
	64	No	TO	<b>4.112</b>	TO	7.065
Floyd	8	No	3.986	3.715	<b>0.232</b>	0.249
vs	16	No	7.171	22.413	<b>1.582</b>	2.108
Processes	32	No	112.793	157.188	<b>4.665</b>	5.619
	64	No	763.232	TO	<b>10.941</b>	32.261
GaussElim	8	No	0.223	0.235	<b>0.177</b>	0.185
vs	16	No	0.599	0.618	0.773	<b>0.481</b>
Processes	32	No	4.398	<b>1.297</b>	4.097	4.155
	64	No	5.748	6.085	5.612	<b>5.429</b>
Heat	8	No	0.323	0.325	0.241	<b>0.225</b>
vs	16	No	<b>0.726</b>	0.789	1.646	1.549
Processes	32	No	4.515	2.497	4.119	<b>1.831</b>
	64	No	4.139	<b>3.706</b>	5.341	9.121
HeatErrors	8	Yes	0.408	0.323	<b>0.221</b>	0.229
vs	16	Yes	1.226	1.772	0.587	<b>0.461</b>
Processes	32	Yes	4.351	2.515	4.568	<b>1.987</b>
	64	Yes	9.492	6.634	6.548	<b>2.932</b>
Integrate	8	No	0.281	<b>0.042</b>	0.251	0.079
vs	16	No	TO	<b>0.223</b>	TO	0.279
Processes	32	No	TO	<b>3.592</b>	TO	3.961
	64	No	TO	<b>4.237</b>	TO	9.119

Diffusion	4	Yes	<b>0.029</b>	0.032	0.115	0.137
(Timesteps=1)	8	Yes	0.037	<b>0.035</b>	0.132	0.125
vs	16	Yes	0.229	<b>0.197</b>	0.295	0.337
Processes	24	Yes	1.457	1.434	<b>0.548</b>	1.482
Diffusion	2	Yes	<b>0.028</b>	0.034	0.131	0.127
(Grid=2x2)	4	Yes	0.031	<b>0.028</b>	0.132	0.127
vs	8	Yes	<b>0.029</b>	0.035	0.121	0.112
Timesteps	16	Yes	<b>0.032</b>	0.038	0.124	0.139
MatMul	8	Yes	0.129	<b>0.069</b>	0.087	0.157
(N=L=M=8)	16	No	4.088	1.269	1.339	<b>0.242</b>
vs	32	No	5.097	<b>1.488</b>	3.714	3.582
Processes	64	No	7.544	4.065	<b>3.835</b>	4.209
MatMul	8	Yes	0.134	<b>0.073</b>	0.087	0.157
(N=L=M=p)	16	Yes	0.869	<b>0.281</b>	0.714	0.356
vs	32	Yes	4.649	<b>0.633</b>	9.545	1.169
Processes	64	Yes	24.643	<b>4.586</b>	98.515	9.554
MatMul	4	No	0.059	<b>0.046</b>	0.068	0.075
(p=8)	6	No	0.092	<b>0.048</b>	0.104	0.067
vs	8	Yes	0.127	<b>0.071</b>	0.098	0.157
Size (N=L=M)	12	Yes	0.189	<b>0.099</b>	0.121	0.187

Table 3: Communication Structure Summaries

Name vs X	X	Trace Size	Epochs			Symmetry	
			(Size, Symmetry, Repeats) ...	Total	Unique	Repeated	Total
Adder vs Processes	8	28	(14,6,1) (2,1,7)	8	2	6	7
	16	60	(30,14,1) (2,1,15)	16	2	14	15
	32	124	(62,30,1) (2,1,31)	32	2	30	31
	64	252	(126,62,1) (2,1,63)	64	2	62	63
Floyd vs Processes	8	176	(20,6,1) (147,0,1) (1,0,9)	11	3	8	6
	16	368	(20,6,1) (23,7,8) (147,0,1) (1,0,17)	27	4	23	13
	32	752	(20,6,1) (23,7,24) (147,0,1) (1,0,33)	59	4	55	13
	64	1520	(20,6,1) (23,7,56) (147,0,1) (1,0,65)	123	4	119	13
GaussElim vs Processes	8	84	(4,1,6) (2,1,2) (1,0,14)	22	3	19	2
	16	172	(4,1,14) (2,1,2) (1,0,22)	38	3	35	2
	32	348	(4,1,30) (2,1,2) (1,0,38)	70	3	67	2
	64	700	(4,1,62) (2,1,2) (1,0,70)	134	3	131	2
Heat vs Processes	8	144	(2,1,60) (1,0,24)	84	2	82	1
	16	296	(2,1,124) (1,0,48)	172	2	170	1
	32	600	(2,1,252) (1,0,96)	348	2	346	1
	64	1208	(2,1,508) (1,0,192)	700	2	698	1
HeatErrors vs Processes	8	144	(17,0,1) (2,1,31) (1,0,15)	47	3	44	1
	16	296	(33,0,1) (2,1,63) (1,0,31)	95	3	92	1
	32	600	(65,0,1) (2,1,127) (1,0,63)	191	3	188	1
	64	1208	(129,0,1) (2,1,255) (1,0,127)	383	3	380	1
Integrate vs Processes	8	28	(14,6,1) (2,1,7)	8	2	6	7
	16	60	(30,14,1) (2,1,15)	16	2	14	15
	32	124	(62,30,1) (2,1,31)	32	2	30	31
	64	252	(126,62,1) (2,1,63)	64	2	62	63

Diffusion (Timesteps=1) vs Processes	4	88	(2,1,16) (18,8,2) (1,0,5)	23	3	20	9
	8	188	(2,1,32) (42,20,2) (1,0,5)	39	3	36	21
	16	388	(90,44,1) (2,1,64) (90,44,1) (1,0,5)	71	4	67	89
	24	588	(138,68,1) (2,1,96) (138,68,1) (1,0,5)	103	4	99	137
Diffusion (Grid=2x2) vs Timesteps	2	150	(2,1,32) (18,8,3) (1,0,8)	43	3	40	9
	4	274	(2,1,64) (18,8,5) (1,0,14)	83	3	80	9
	8	522	(2,1,128) (18,8,9) (1,0,26)	163	3	160	9
	16	1018	(2,1,256) (18,8,17) (1,0,50)	323	3	320	9
MatMul (N=L=M=8) vs Processes	8	54	(46,5,1) (1,0,8)	9	2	7	5
	16	78	(2,1,7) (48,7,1) (1,0,16)	24	3	21	8
	32	126	(2,1,23) (48,7,1) (1,0,32)	56	3	53	8
	64	222	(2,1,55) (48,7,1) (1,0,64)	120	3	117	8
MatMul (N=L=M=p) vs Processes	8	54	(46,5,1) (1,0,8)	9	2	7	5
	16	110	(94,13,1) (1,0,16)	17	2	15	13
	32	222	(190,29,1) (1,0,32)	33	2	31	29
	64	446	(382,61,1) (1,0,64)	65	2	63	61
MatMul (p=8) vs Size (N=L=M)	4	38	(2,1,3) (24,3,1) (1,0,8)	12	3	9	4
	6	46	(2,1,1) (36,5,1) (1,0,8)	10	3	7	6
	8	54	(46,5,1) (1,0,8)	9	2	7	5
	12	70	(62,5,1) (1,0,8)	9	2	7	5

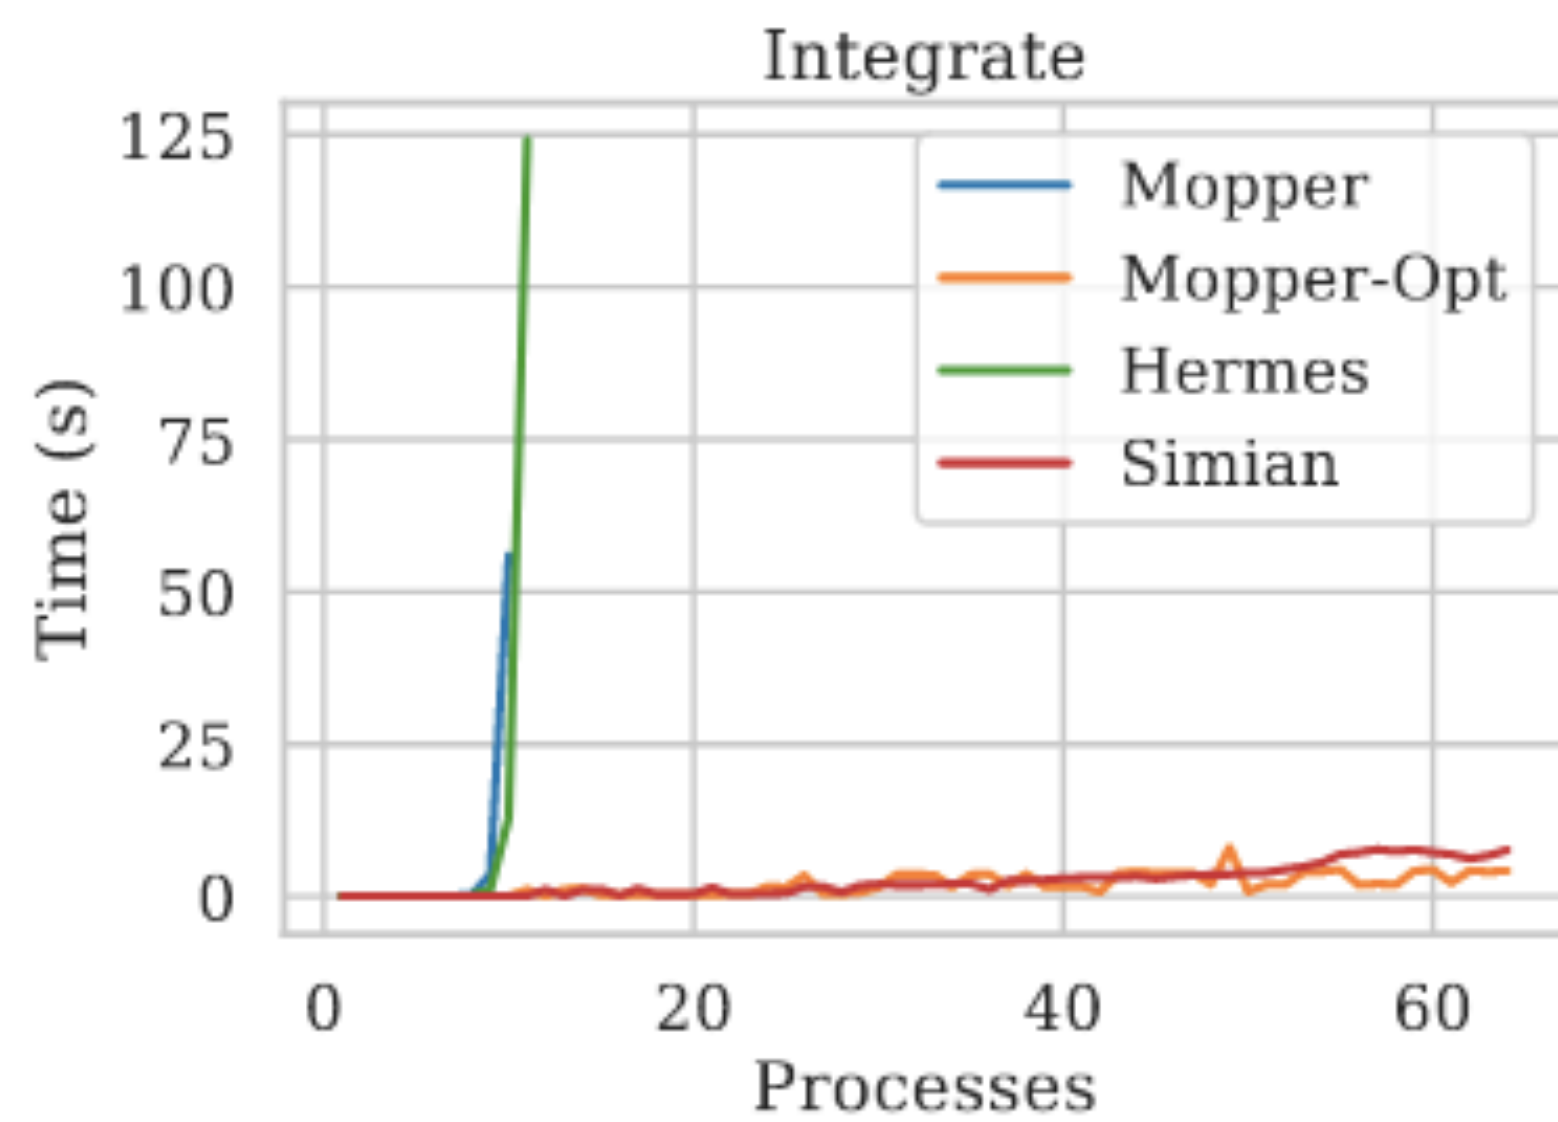


**Table 4: Component-wise Times (in s)**


Name vs X	X	Time		
		Scheduler	Analyzer	Solver
Adder	8	<b>0.031</b>	0.014	0.013
vs	16	<b>1.084</b>	0.035	0.036
Processes	32	<b>1.381</b>	0.206	0.238
	64	2.358	1.957	<b>2.938</b>
Floyd	8	0.042	0.042	<b>0.543</b>
vs	16	<b>1.290</b>	0.094	0.553
Processes	32	<b>1.708</b>	0.239	0.541
	64	<b>4.764</b>	0.753	0.558
GaussElim	8	<b>0.151</b>	0.022	0.014
vs	16	<b>0.223</b>	0.032	0.014
Processes	32	<b>2.273</b>	0.028	0.011
	64	<b>3.165</b>	0.045	0.011
Heat	8	<b>0.264</b>	0.050	0.009
vs	16	<b>1.468</b>	0.031	0.006
Processes	32	<b>4.195</b>	0.052	0.008
	64	<b>3.119</b>	0.099	0.007
HeatErrors	8	<b>0.353</b>	0	0
vs	16	<b>0.565</b>	0	0
Processes	32	<b>2.408</b>	0.073	0.016
	64	<b>3.407</b>	0.183	0.019
Integrate	8	<b>0.032</b>	0.015	0.013
vs	16	<b>0.058</b>	0.035	0.037
Processes	32	<b>1.372</b>	0.198	0.235
	64	2.682	1.972	<b>2.925</b>

Diffusion	4	0.039	<b>0.061</b>	0.022
(Timesteps=1)	8	0.038	<b>0.585</b>	0.091
vs	16	0.135	<b>6.635</b>	1.783
Processes	24	1.021	<b>27.953</b>	7.294
Diffusion	2	0.022	<b>0.185</b>	0.021
(Grid=2x2)	4	0.027	<b>0.968</b>	0.021
vs	8	0.106	<b>6.952</b>	0.021
Timesteps	16	0.982	<b>65.333</b>	0.021
MatMul	8	<b>0.032</b>	0.018	0.026
(N=L=M=8)	16	<b>0.065</b>	0.024	0.023
vs	32	<b>1.389</b>	0.026	0.024
Processes	64	<b>2.618</b>	0.032	0.022
MatMul	8	<b>0.032</b>	0.017	0.026
(N=L=M=p)	16	0.066	0.043	<b>0.112</b>
vs	32	<b>1.699</b>	0.232	0.834
Processes	64	2.331	2.109	<b>11.731</b>
MatMul	4	<b>0.032</b>	0.017	0.015
(p=8)	6	<b>0.031</b>	0.018	0.021
vs	8	<b>0.032</b>	0.017	0.026
Size (N=L=M)	12	0.034	0.028	<b>4.875</b>

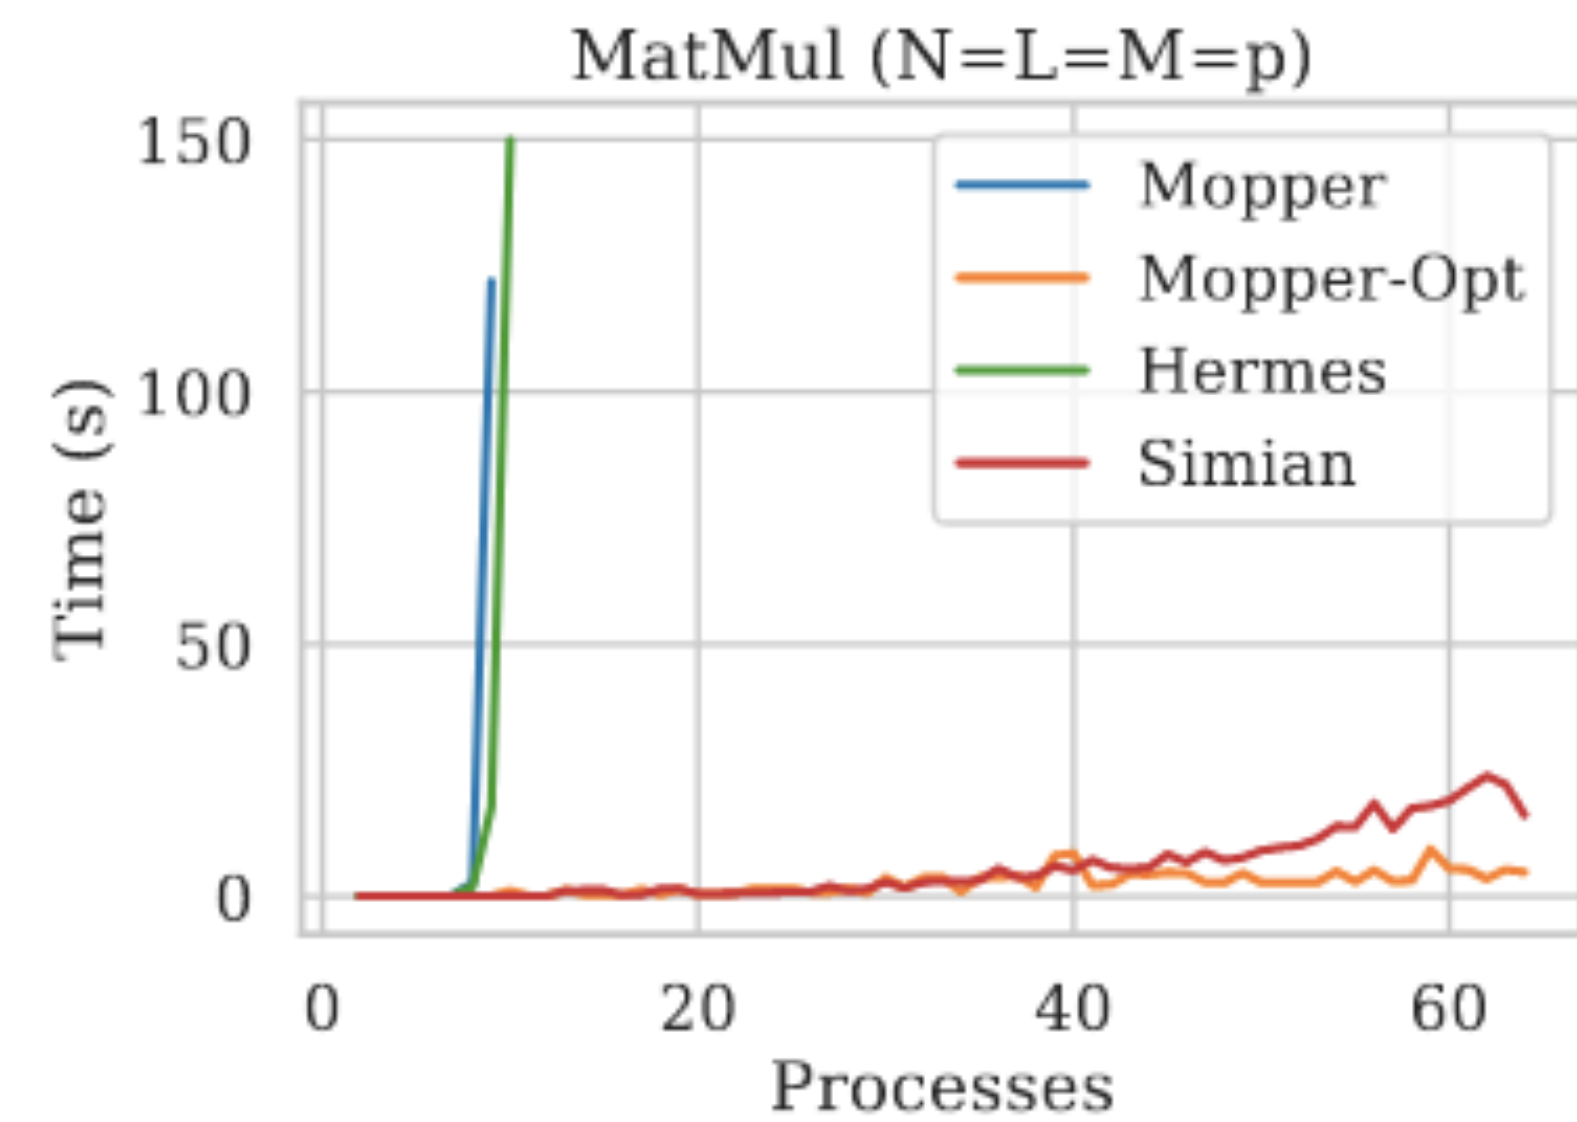




(a) Integrate




(b) Diffusion (Timesteps=1)

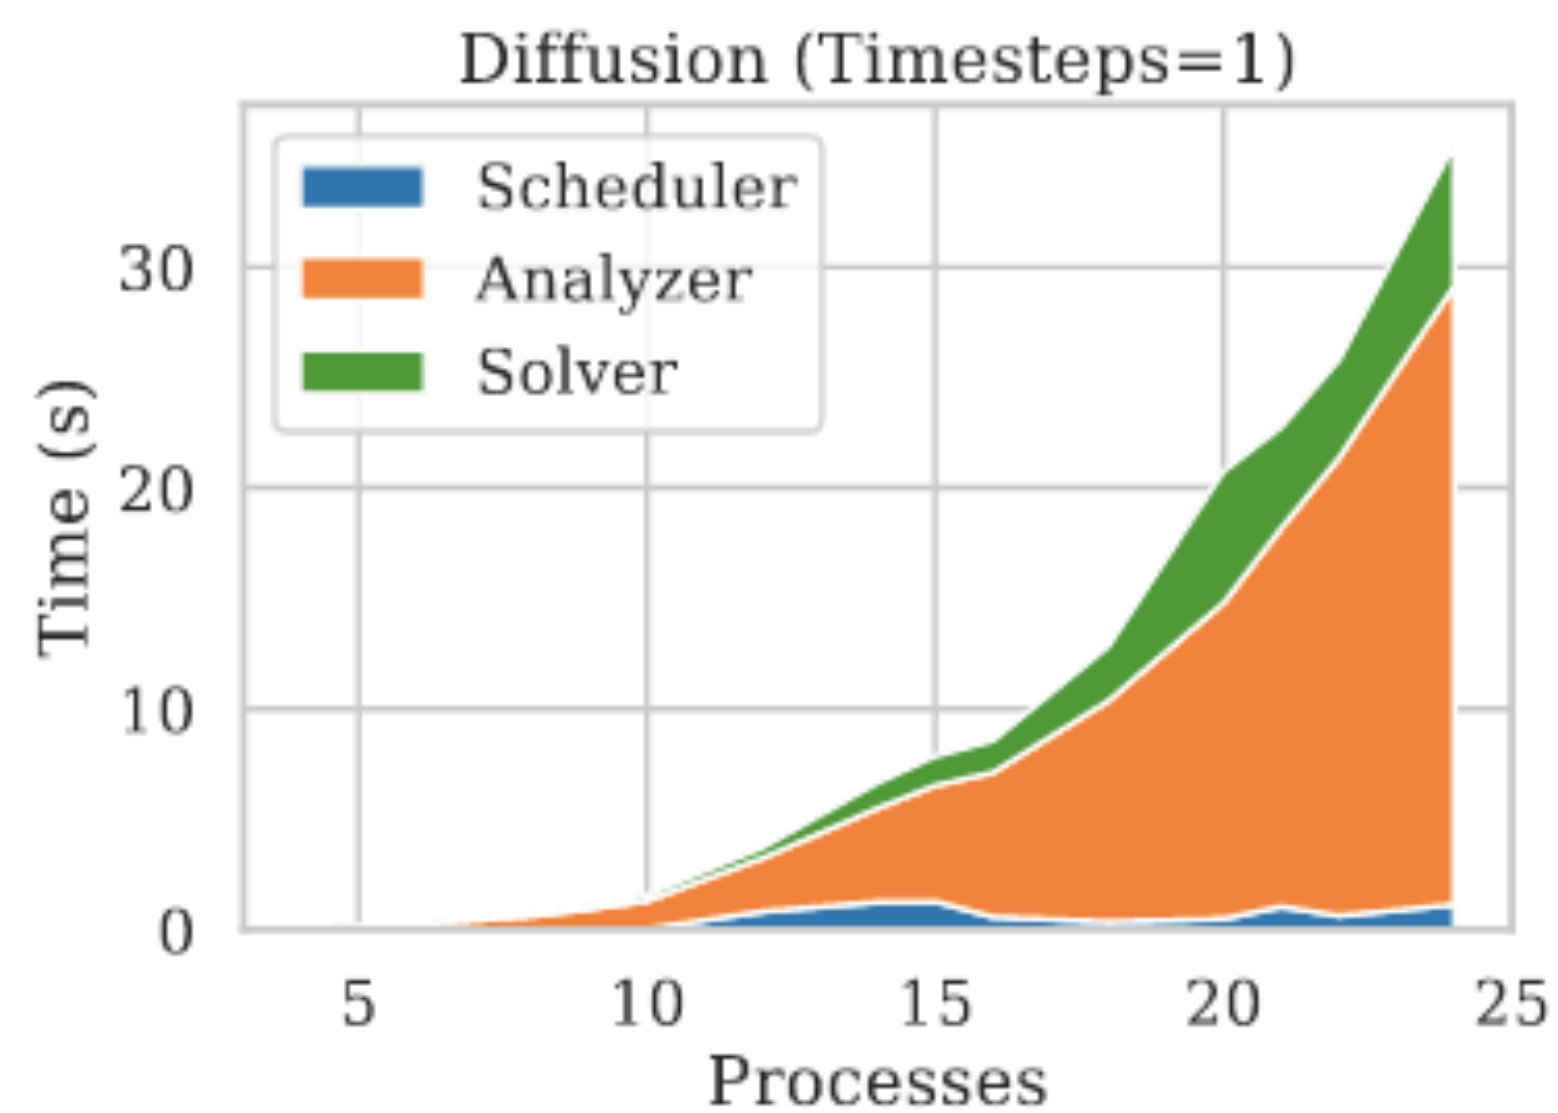


(c) MatMul (N=L=M=p)

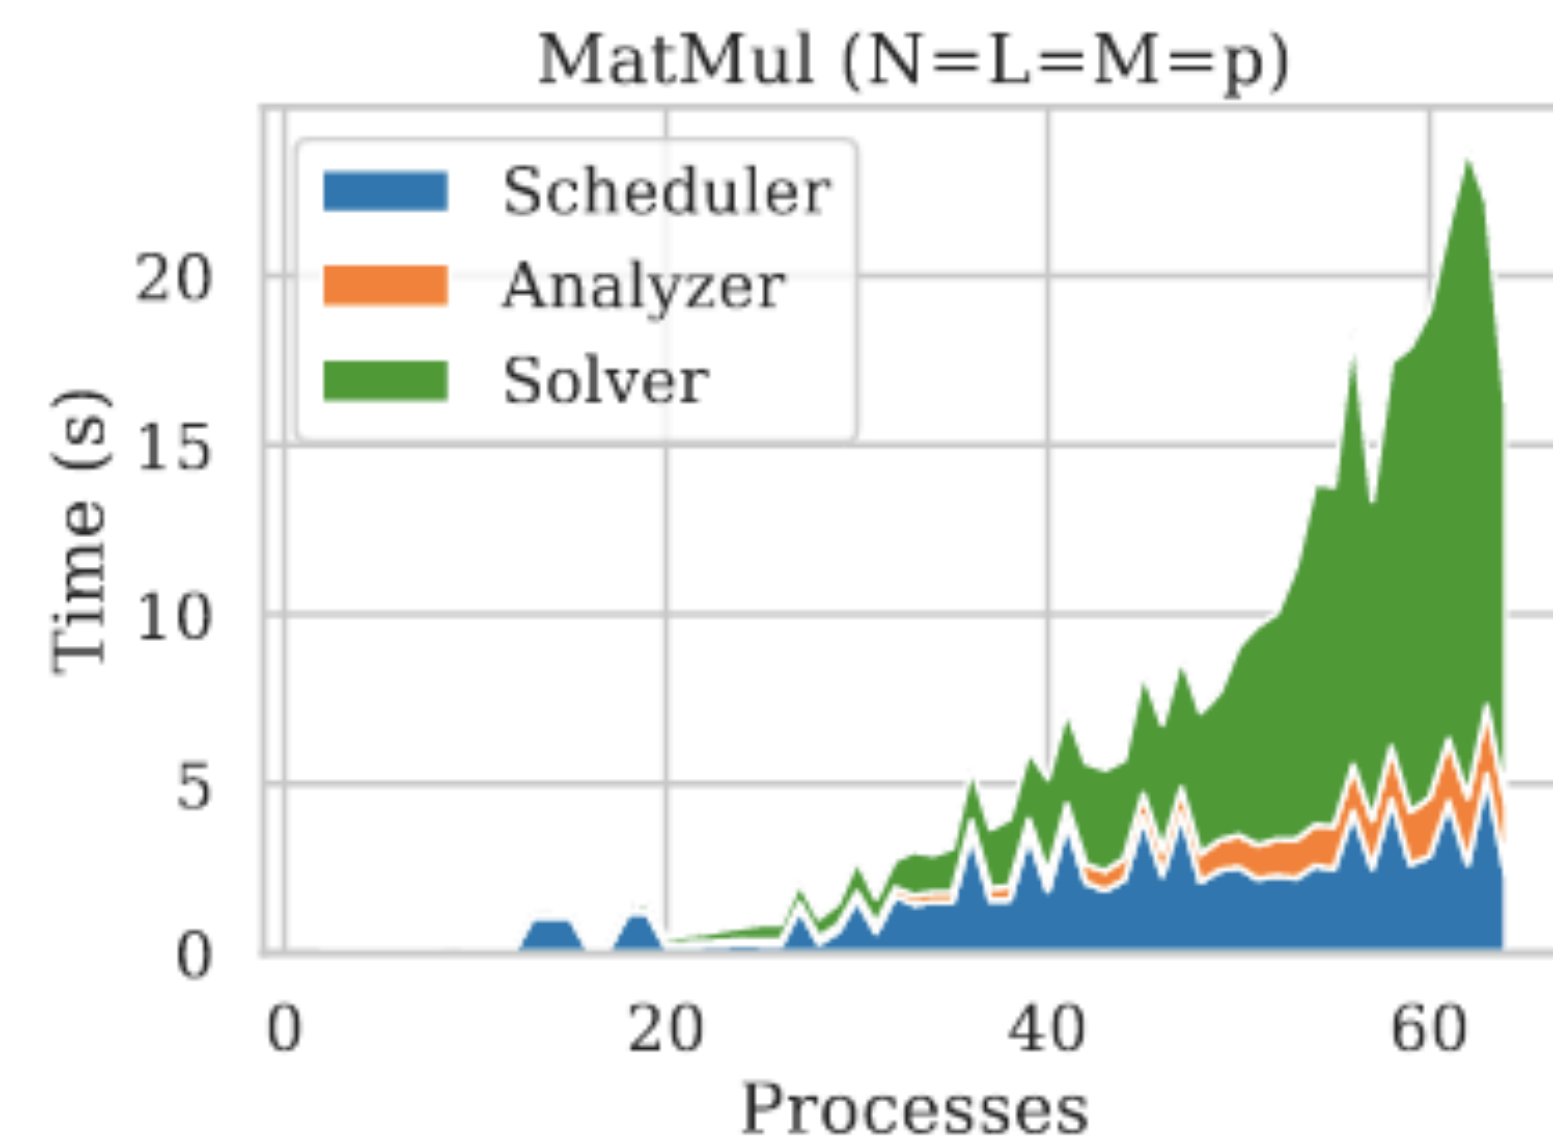
Figure 3: Total Times



(a) Integrate



(b) Diffusion (Timesteps=1)



(c) MatMul (N=L=M=p)

**Figure 4: Component Times**

# Future Directions

- Extension to **multi-path programs** where communication affects control flow
- Applicability of our techniques to **other concurrency frameworks**
- Expanding the **scope of our tool** and evaluating against recent benchmarks

**Thank you!**  
**Questions?**